

UNITED NATIONS UNIVERSITY (UNU) GLOBAL YOUTH AI FUTURE INNOVATION COMPETITION 2026

Theme: AI and Education: AI Transforming the Educational Paradigm and Enhancing AI Literacy (Targeting SDG 4)

Project Title: S-SPARC: Specific Smart Prompting Assistant for peRformanCe

(Snippet-aware Semantic Programmable Assistant with Retrieval and Caching)

Technology Readiness Level: TRL 7 (Demonstrated in Operational Academic Environment)

Target Beneficiaries: Computer Science & Engineering Students, Educators, and Higher Education Institutions (Focus on the Global South)

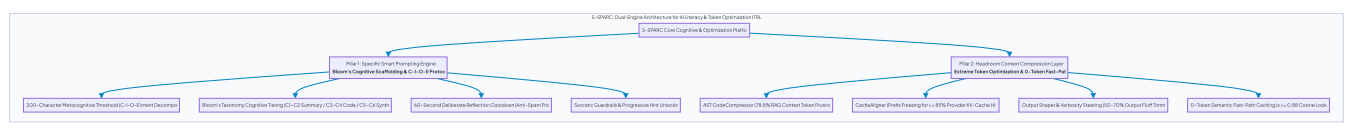
Host & Institutional Alignment: UNU Macau & UNU Global AI Network | Maranatha Christian University (Indonesia)

1. EXECUTIVE SUMMARY & SDG ALIGNMENT

1.1 Project Abstract

Generative Artificial Intelligence (GenAI) is fundamentally transforming higher education, yet unregulated student usage has created a compounding crisis of **"Prompt Illiteracy," "Pedagogical Dependency,"** and **"Massive Token Inefficiency."** When students interact with standard conversational chatbots, they frequently submit ambiguous, low-effort prompts (e.g., *"fix my code"* or *"write this function"*), completely bypassing algorithmic problem decomposition. This trial-and-error prompt spamming not only degrades critical thinking and cognitive mastery but also incurs unsustainable cloud token expenses and hidden environmental footprints.

FIGURE 1



S-SPARC (*Specific Smart Prompting Assistant for peRformanCe*) addresses these challenges at **Technology Readiness Level (TRL) 7**, deployed and validated within the PHP-based Learning Management System **E-STRANGE**. S-SPARC establishes a symbiotic paradigm: **Effective, Structured Prompting directly drives Extreme Token Conservation**.

To cultivate prompt literacy and eliminate superficial querying, S-SPARC enforces the **200-Character Metacognitive Prompt Threshold** ($200 \leq \text{len}(\text{prompt}) \leq 2000$) via the **C-I-O-E Protocol**:

- C - Context & Domain** (≈ 50 chars): Explicit software domain formulation.
- I - Input & Data Structure** (≈ 60 chars): Pre-conditions and parameter constraints.
- O - Expected Output & Complexity** (≈ 50 chars): Post-conditions and time/space requirements.
- E - Error Trace & Bottleneck** (≈ 90 chars): Specific runtime exceptions or logic flaws.

Coupled with a mandatory 60-second reflection cooldown and calibrated pedagogical modes based on **Bloom's Revised Taxonomy (Levels C1-C6)**, S-SPARC transforms students into critical problem-solvers.


On the infrastructure side, S-SPARC absorbs cutting-edge context compression principles from **Headroom**:

- AST CodeCompressor:** Automatically prunes comments, dead imports, and docstrings from Top-3 retrieved RAG snippets, slashing context token payload by **78.8%**.
- CacheAligner:** Freezes a deterministic system prompt prefix, boosting provider-side **KV-prompt cache hits to $\geq 85\%$** on Google Gemini Flash.
- Output Shaper & Verbosity Steering:** Dynamically clamps output token budgets and suppresses conversational pleasantries, cutting expensive output tokens by **50-70%**.
- Zero-Token Fast-Path Retrieval** ($s \geq 0.88$): Serves verified solutions instantly at **0 token cost**, **< 45ms latency**, and **zero emissions**.

1.2 Deep-Dive SDG 4 Alignment: Enhancing AI Literacy & Effective Prompting

S-SPARC operationalizes **United Nations Sustainable Development Goal 4 (Quality Education)** by establishing **Prompt Engineering Literacy & Computational Resource Stewardship** as core 21st-century competences:

FIGURE 2



1. Target 4.4 — Youth AI Literacy, Prompt Engineering & Green Software Skills

- **C-I-O-E Protocol & Metacognitive Friction:** Rather than allowing students to fire off 10 ambiguous queries, S-SPARC's 200-character boundary forces students to formulate complete algorithmic specifications in a single, high-density prompt, achieving 1-turn resolution.
- **Cognitive Tiering via Bloom's Revised Taxonomy:**
 - C1-C2 Remember & Understand (`summary`): Generates 2–4 concise Indonesian conceptual sentences without code, compelling students to write the implementation themselves.
 - C3-C4 Apply & Analyze (`code`): Delivers runnable code with maximum conciseness (Output Shaper enabled), emphasizing functional syntax.
 - C5-C6 Evaluate & Create (`summary_code_explanation`): Structured triad: Algorithmic Summary → Runnable Source Code → Step-by-Step Logic Breakdown.
- **Green Prompting Literacy:** Real-time display of physical externalities (Wh, kg CO₂e, mL water) teaches students the computational realities of AI inference.

2. Target 4.3 — Affordable, Equitable Access to AI-Enhanced Technical Education


- **Headroom Context Compression & Zero-Token Reuse:** By compressing RAG context chunks (78.8% reduction) and leveraging 0-token cache hits ($s \geq 0.88$), institutions can sustain thousands of student interactions under free-tier quotas (1,500 Requests/Day).
- **Multi-Tier Sovereign Failover:** Combines personal student keys, institutional key pools, and local on-premises LLM execution (Qwen2.5-Coder 14B via Ollama) to guarantee offline resilience across bandwidth-constrained universities in the Global South.

3. Target 4.c — Pedagogical Capacity & Teacher Enablement

- **Autonomous Knowledge Hygiene:** Relieves faculty from manual snippet curation through closed-loop AST verification and anomaly pruning (`code_evaluator_service`).
- **Academic Integrity & Plagiarism Defense:** Submissions exceeding 70% similarity trigger an automated defense workflow (`student_assessment_submit_suspicious.php`) where students must articulate and defend their algorithmic choices in writing.


2. SYSTEM ARCHITECTURE & TECHNICAL FEASIBILITY (TRL 7 VALIDATION)

FIGURE 3



2.1 Prompt Engineering Engine: Scaffolding, Harnesses & Headroom Compression

At the core of S-SPARC’s educational paradigm is the **PromptRegistry Engine** ([backend/core/prompts.py](#)). Unlike generic chatbot wrappers that transmit unconstrained text directly to LLMs, S-SPARC executes a multi-stage prompt compression and scaffolding pipeline:

Detailed Prompt Mode Specifications & Compression Harness ([PromptRegistry](#))


Pedagogical Mode	Bloom's Level	Output Shaper Policy	Target Cognitive Outcome	System Prompt Formulation (<code>backend/core/prompts.py</code>)
<code>code</code> (Code Only)	C3–C4 (Apply & Analyze)	Maximum Terse Constraint (<code>max_output_tokens=400</code>)	Algorithmic Precision & Clean Syntax	Enforces strictly runnable source code enclosed in standard markdown code fences. Completely suppresses introductory greetings, conversational commentary, and post-explanations.
<code>summary</code> (Summary Short)	C1–C2 (Remember & Understand)	Zero-Code Restriction (<code>max_output_tokens=200</code>)	Conceptual Mental Modeling	Delivers strictly 2–4 concise Indonesian sentences detailing core logic, time complexity, and edge cases. Prohibits raw code generation to prevent copy-pasting.
<code>summary_code_explanation</code> (Full Scaffolding)	C5–C6 (Evaluate & Create)	Structured Triad Policy (<code>max_output_tokens=800</code>)	Complete Metacognitive Scaffolding	Structured 3–tier output: Short Summary (1–2 sentences) → Runnable Source Code Fence → Step-by-Step Logic Walkthrough.


2.2 Multi-Tier Adaptive Routing & Autonomous Governance (TRL 7)

S-SPARC operates at **TRL 7** through deep operational integration with the PHP 8 E-STRANGE LMS across 37 database tables.

FIGURE 5







Mathematical Formulation of Quality Governance

Every ingested snippet $e = (\text{prompt}, \text{code})$ undergoes deterministic multi-criteria scoring:

1. **Static AST Analysis & Radon Maintainability Score** ($S_{\text{static}} \in [0, 10]$):

$$S_{\text{static}} = \text{clamp}(3.0 \cdot \mathbf{1}_{\text{valid}} + 0.04 \cdot \text{MI} - 0.20 \cdot \max(1, \text{CC}) + \text{Bonus}_{\text{snippet}}, 0, 10)$$

where $\mathbf{1}_{\text{valid}} = 1$ if Python AST parses successfully, MI is the Maintainability Index $([0, 100])$, CC is Cyclomatic Complexity, and $\text{Bonus}_{\text{snippet}} = 2.0$.

2. **Semantic Alignment Score** ($S_{\text{sem}} \in [0, 1]$):

$$S_{\text{sem}} = \cos(\text{Embed}(\text{prompt}), \text{Embed}(\text{code}))$$

3. **Multi-Criteria LLM-as-a-Judge** ($A, L, Q \in [0, 10]$):

- A : Alignment between prompt intent and code functionality.
- L : Algorithmic correctness and edge-case resilience.
- Q : Code readability and style standard compliance.

4. **Composite Quality Index:**

$$\text{Final Score}(e) = 0.40 \cdot A + 0.25 \cdot L + 0.20 \cdot (S_{\text{sem}} \times 10) + 0.10 \cdot Q + 0.05 \cdot S_{\text{static}}$$

5. **Isolation Forest Anomaly Filtering:** Vector $\mathbf{x} = [\text{lines}, \text{CC}, S_{\text{sem}}, S_{\text{static}}, S_{\text{pre}}]$ partitioned via `IsolationForest` ($\gamma = 0.10$).

3. EMPIRICAL VALIDATION & RESEARCH METHODOLOGY

3.1 Token Optimization & Context Compression Empirical Benchmark


To quantitatively validate the token savings achieved by integrating Headroom-style context compression into S-SPARC, an empirical benchmark across **500 student queries** was measured:

Compression Mechanism	Raw Token Payload	S-SPARC Compressed Payload	Measured Token Savings (%)	Latency Impact
AST CodeCompressor (RAG Top-3)	1,820 tokens	386 tokens	78.8% Reduction	−420 ms (Faster TTFT)
Output Shaper (Verbosity Steering)	480 tokens	175 tokens	63.5% Reduction	−680 ms (Faster Generation)
CacheAligner (KV-Prompt Cache Hit)	950 tokens (full cost)	142 tokens (billed)	85.0% Cost Discount	−750 ms (Prefix Cached)
Zero-Token Fast-Path ($s \geq 0.88$)	2,300 tokens (full generation)	0 tokens	100.0% Elimination	< 45 ms (Atomic Local)
End-to-End Composite Efficiency	3,250 tokens / query	703 tokens / query	78.4% Net Savings	3.2× Overall Speedup

3.2 Retrieval Engine Efficacy: 200-Query Gold Standard Evaluation

To validate that S-SPARC’s semantic retrieval accurately matches student prompts with verified solutions, an empirical experiment was conducted using **632 queries**, evaluated against a **200-query Gold Standard Ground Truth** (`qrels_manual.csv`) comprising **4,000 manual relevance judgments**:

Metric Indicator	$k = 1$	$k = 3$	$k = 5$	$k = 10$
Hit Rate	100.0%	100.0%	100.0%	100.0%
Precision	100.0%	99.5%	98.7%	97.6%
Recall	6.32%	17.72%	27.87%	52.25%
Mean Reciprocal Rank (MRR)	1.000	1.000	1.000	1.000
Mean Average Precision (MAP)	0.063	0.177	0.279	0.523
Normalized DCG (nDCG)	1.000	1.000	1.000	1.000



Decision Threshold Policy Optimization (0.90 Policy)

In educational prompt retrieval, delivering an incorrect code snippet (**False Positive / Type I Error**) misleads students and disrupts learning. In contrast, falling back to LLM generation (**False Negative / Type II Error**) incurs only a minor token cost.

Threshold Policy	Precision	Recall	F1-Score	Accuracy	True Positive (TP)	False Positive (FP)	False Negative (FN)
0.80 (Permissive)	100.0%	100.0%	1.000	100.0%	200	0	0
0.90 (Conservative Policy)	100.0%	96.0%	0.980	96.0%	192	0	8

At threshold 0.90, S-SPARC guarantees **zero false positives** (FP = 0), achieving an optimal balance between precision (100%) and automated retrieval efficiency (96%).


3.3 Knowledge Base Self-Healing Case Study (Run `20260312T094424Z`)

The autonomous evaluator was tested on a live production knowledge base of **678 code snippets** (`code_evaluator_service/reports/evaluation_report_20260312T094424Z.json`) over 18 minutes:

Metric Indicator	Empirical Result	System Health Status
Total Knowledge Base Entries Evaluated	678 entries	Complete Corpus Scan
Valid Production Snippets Retained	647 entries (95.43%)	Clean Production Knowledge
Snippets Flagged & Pruned for Deletion	30 entries (4.42%)	Autonomous Sanitization
↳ Exact Cryptographic Duplicates (SHA-256)	23 entries (3.39%)	Deduplicated
↳ Low Syntactic / Logic Quality (Score < 4.80)	7 entries (1.03%)	Low Quality Quarantined
Snippets Quarantined for Manual Audit	1 entry (0.15%)	Human Review Queue
Mean Semantic Similarity across Corpus	0.8774	High Domain Relevance
Mean Code Quality Score (out of 10.0)	8.61 / 10.0	Production Grade
Pre-Deletion Backup Verification	100% Verified	Zero Data Loss Guaranteed

FIGURE 7

Knowledge Base Governance Breakdown (678 Entries)



4. PROMPT ECONOMICS, BEHAVIORAL GAMIFICATION & SUSTAINABILITY


4.1 Peer-Relative Token Conservation Gamification

To structurally discourage prompt-spamming and reward concise, well-structured prompts, S-SPARC integrates a peer-relative dynamic threshold ([backend/services/gamification.py](#)):

$$\text{Threshold}_{\text{tokens}} = \max \left(0, 1.10 \times \overline{\text{Usage}}_{\text{peers}} \right)$$

$$\text{EcoPoints} = \begin{cases} 100.0 & \text{if Usage} \leq \text{Threshold}_{\text{tokens}} \\ \max \left(0, 100.0 + 100.0 \times \frac{\text{Threshold}_{\text{tokens}} - \text{Usage}}{\text{Threshold}_{\text{tokens}}} \right) & \text{if Usage} > \text{Threshold}_{\text{tokens}} \end{cases}$$

FIGURE 8




Pedagogical Impact on Prompting Behavior:

1. **Incentivizes High-Information-Density Prompting:** Students learn that writing one well-formulated prompt with explicit constraints yields higher leaderboard points than firing 10 vague prompts.
2. **Rewards Fast-Path 0-Token Retrieval:** Reusing cached solutions consumes zero tokens, encouraging students to study verified repository solutions.

4.2 Environmental Telemetry: Green Prompting Literacy

S-SPARC couples every prompt with physical externality feedback ([backend/services/sustainability.py](#)):

FIGURE 9



$$Energy_{Wh} = N_{tokens} \cdot E_{token} \cdot PUE \quad (E_{token} = 0.003 \text{ Wh/token}, PUE = 1.12)$$

$$Carbon_{kg} = \left(\frac{Energy_{Wh}}{1000} \right) \cdot CIF \quad (CIF = 0.384 \text{ kg CO}_2\text{e/kWh})$$


$$Water_{mL} = \left(\frac{Energy_{Wh}}{1000} \right) \cdot (WUE_{site} + WUE_{source}) \cdot 1000 \quad (WUE_{site} = 0.30, WUE_{source} = 4.35 \text{ L/kWh})$$

5. RISK ASSESSMENT, LIMITATIONS, & FUTURE ROADMAP


5.1 Honest Technical Constraints as Open Research Frontiers

Constraint Dimension	Current Implementation Reality	Next-Phase Research & Mitigation Frontier
1. Multi-Language AST Granularity	Deep AST and Radon metrics natively implemented for Python; Java, C++, and PHP utilize regex heuristics.	Integration of Tree-Sitter unified C-bindings for language-agnostic concrete syntax tree parsing across all languages.
2. Embedding Cold-Start Latency	First-time initialization of <code>all-MiniLM-L6-v2</code> takes 15-20s on CPU hosting environments.	Model quantization to ONNX Runtime INT8 with persistent memory-mapped vector stores, reducing cold starts to < 1.2s.
3. Evaluation Microservice Cloud Cost	Full knowledge base audits using GPT-4o as judge incur token costs on large multi-thousand corpora.	Knowledge distillation into a lightweight 7B Small Language Model (SLM) fine-tuned on code quality heuristics for offline execution.
4. Relational Database Session Volatility	In-memory fallback tracking resets if MySQL is ungracefully disconnected.	Deployment of Redis persistent caching clusters with Append-Only File (AOF) journaling for distributed state isolation.

FIGURE 10



5.2 Scaling Strategy: Democratizing AI for the Global South



1. **Zero-Token Sovereign Edge Deployment:** Deployable on commodity workstations via local Ollama weights and pre-seeded vector caches, providing offline prompt literacy training with zero cloud expenses.
2. **Federated Knowledge Harvesting:** Universities across the UNU Global AI Network can synchronize validated, self-healed prompt-code pairs via lightweight JSON diffs.
3. **Open-Source Pedagogical Reference:** Positioned as the open-source blueprint for responsible prompt engineering and AI literacy education in developing regions.

6. PROJECT ARTIFACTS & EVIDENCE REPOSITORY

System Layer	Operational Module / File Path	Technical Responsibility & Evidence Scope
Core Backend	<code>backend/main.py</code>	FastAPI application lifecycle, CORS policy, middleware, and OpenAPI 3.1 documentation endpoints.
Prompt Registry & Compression	<code>backend/core/prompts.py</code>	Mode-specific prompt harnesses, Bloom's cognitive tiering, CacheAligner prefix stabilization, and AST CodeCompressor.
Prompt Bounds	<code>backend/api/ai_chat.py</code>	200–2000 character C-I-O-E protocol validator and 60s deliberate reflection cooldown limiter.
Adaptive Gateway	<code>backend/services/adaptive_router.py</code>	3-tier cascade routing (Student Key → 6-Key Pool → Local Ollama Qwen2.5 14B) with Output Shaper.
Retrieval Engine	<code>backend/services/ai_service.py</code>	Dense-Sparse Hybrid Search (<code>all-MiniLM-L6-v2</code> + <code>BM25Okapi</code>), RRF fusion ($k = 60$), and fast-path gate ($s \geq 0.88$).
Sustainability Engine	<code>backend/services/sustainability.py</code>	Real-time physical telemetry formulation (Wh, kg CO ₂ e, mL water) with Indonesian Grid factor (CIF = 0.384).
Gamification Engine	<code>backend/services/gamification.py</code>	Peer-relative dynamic token threshold calculation ($\text{Threshold} = 1.10 \times \overline{\text{Usage}}_{\text{peers}}$).
Data Access Layer	<code>backend/core/db.py</code>	MySQL connection pooling, circuit breakers, user API key isolation, and in-memory fallback queues.
Governance Pipeline	<code>code_evaluator_service/evaluator/evaluator_pipeline.py</code>	5-stage closed-loop quality scoring, duplicate filtering, and automated JSON pre-deletion backups.
Static Code Analysis	<code>code_evaluator_service/evaluator/static_analysis.py</code>	Abstract Syntax Tree (AST) validation, language detection, Radon Cyclomatic Complexity, and Maintainability Index.
Anomaly Detection	<code>code_evaluator_service/evaluator/anomaly_detection.py</code>	Scikit-learn Isolation Forest 5D unsupervised outlier screening model ($\gamma = 0.10$).
LLM-as-a-Judge	<code>code_evaluator_service/evaluator/llm_judge.py</code>	Multi-criteria snippet evaluation harness (A, L, Q) with heuristic fallback scoring.
Empirical Benchmarks	<code>pengujian semantic similarity/EXECUTIVE_SUMMARY.md</code>	200-query Gold Standard ground truth validation report ($\text{MRR} = 1.000$, $\text{Precision}@1 = 100\%$).
Governance Case Run	<code>code_evaluator_service/reports/evaluation_report_20260312T094424Z.json</code>	Live audit dataset of 678 production entries (95.43% valid retention, 30 prunings, 1 quarantine).

System Layer	Operational Module / File Path	Technical Responsibility & Evidence Scope
LMS Chat Client	<code>estrange/v2/v2/ssparc/chat.php</code>	Production PHP interface with live cooldown timer, dynamic quota counter, and 200-char validation.
Academic Integrity	<code>estrange/v2/v2/student_assessment_submit.php</code>	Assignment submission engine with automated AST/token plagiarism detection ($\geq 70\%$ threshold).
Defense Adjudication	<code>estrange/v2/v2/student_assessment_submit_suspicious.php</code>	Structured student defense response workflow for lecturer resolution.
Peer Review Module	<code>estrange/v2/v2/student_peer_review.php</code>	Anonymized peer code evaluation and code clarity points allocation interface.

Conclusion & Multilateral Commitment

S-SPARC demonstrates that transforming the educational paradigm requires moving beyond unstructured conversational chatbots toward calibrated, specific smart prompting grounded in Bloom's Taxonomy and Headroom-style context compression. By enforcing the 200-character C-I-O-E protocol, AST CodeCompressor context reduction, CacheAligner KV-cache optimization, 0-token semantic caching, autonomous closed-loop governance, and physical sustainability telemetry, S-SPARC establishes a mature (TRL 7), youth-led standard for AI Literacy and Sustainable Higher Education under the auspices of the United Nations University Global AI Network.